

Full-Stack GeoAI


Building A World On Autopilot (Worth Living In)

Krzysztof Janowicz • University of Vienna

The 17th Conference on Spatial Information Theory (COSIT) 22-25 September, York, UK

[Including COSIT easter eggs and references for you to enjoy.]

Envisioning a Full-Stack GeoAI




Motivating Claims

Core premise*:

Foundation models are increasingly eroding the need for domain/task-specific models.


- The era of *ever more spatially explicit model* design will come to an end.
- As we increasingly build on **foundation** models, understanding them **beyond accuracy** alone will become key.
- GeoAI has historically targeted mostly **levels 0-2**; with great new papers at COSIT 2026.
- Many **newly important (theoretical)** questions will arise at levels 3 & 4.
- Future work must establish a **spatial*** **perspective** across the entire AI cycle from problem framing to deployment effects.



* If not stated otherwise explicitly, spatial will include spatiotemporal and platial.

* Keep in mind, in a chain of arguments you cannot question the premise ;-)

The (Idealized) Jagged Frontier of AI and Its Hidden Assumptions





[Inspired by Dell'Acqua et al. (2023)]

We Misjudge the Frontier

- Many **types** of charts and graphics show the **jagged** AI frontier; most fall victim to the same **hidden assumptions**.
 - **Increasing difficulty** for humans does not imply increasing difficulty for AI
 - Misleading (often linear) **interpolation** between tasks
 - Assuming that performance beyond the human-skill frontier implies **safe automation**
 - Regarding skills and knowledge as isolated, point-like per-task capabilities, while **knowledge is a graph** (van Harmelen, 2011)
 - Assuming a meaningful frontier can be **defined by accuracy** in the first place
 - ...
- Hence, we risk confusing **superhuman abilities** for **superhuman intelligence**.

Task-Completion Time Horizons





4 notions of time, 2 hidden

- 1 2-3 year evaluation period
- 2 Task completion time
- 3 Time the AI takes
- 4 Time to evaluate completion

Figure source: METR 2026; <https://metr.org/time-horizons/>

Task-Completion Time Horizons



4 notions of time, 2 hidden

1 2-3 year evaluation period

2 Task completion time

3 Time the AI takes

- Largely irrelevant
- 10k agents can run in parallel
- Millennium Prize problem example

4 Time to evaluate completion

- *Hard to solve, easy to verify* class of problems
- What if *correct* itself is uncertain, vague, or contested?
- **Claim:** many *geog. problems* do not fall in that class

Figure source: METR 2026; <https://metr.org/time-horizons/>

Three Perspectives on Intelligence

It seems useful to approach computational accounts of intelligence through three* complementary perspectives.

■ **Intelligence as Inference** (*symbolic perspectives*)

A system is intelligent if it can **derive new implicit statements from existing explicit statements**.

$$\text{Waterfall-KB} \models \neg \exists p \left(\text{Falling}(p) \wedge \text{enacts}^-(p, wf) \right)$$

■ **Intelligence as Compression** (*neural perspectives*)

A system is intelligent if it can **compress** (abstract, or generalize over) observations by **recognizing patterns**.

$$P(\text{"falls"} \mid \text{"The 198m tall Cautley Spout"})$$

■ **Intelligence as Invariance** (*today's talk perspective*)

A system is intelligent if it remains **stable under semantics-preserving transformations** yet adapts to semantically relevant **context[†]**.

[†]For the sake of this talk, context will be *geographic*

*Many more, e.g., Microsoft using a purely financial AGI definition

Geographic Bias in Training Data and Beyond

- Open Images and ImageNet only contained 3% of data from China and India combined
- Researchers studying **geo-diversity** utilize key concepts of geographic data analysis, e.g., MAUP.
- **GeoAI** is not an AI4Geo one-way-street; we have **a lot to offer** to the broader community.
- Overcorrective **debiasing** is problematic, e.g., Google's black Nazi soldiers.
- How can we define and measure **geo-diversity**?
- Bias is a great example as it **spans across multiple** full-stack levels
- **We will deliberately take another perspective here by focusing on invariance**





Figure 1: Fraction of Open Images and ImageNet images from each country. In both data sets, top represented locations include the US and Great Britain. Countries are represented by their two-letter ISO country codes.



Figure 2: Distribution of the geographically identifiable images in the Open Images data set, by country. Almost a third of the data in our sample was US-based, and 60% of the data was from the six most represented countries across North America and Europe.

(Fig. by Shankar, 2017)


A Geographic Sustainability Perspective on Tokenization



- Tokenization is often treated as a purely technical pre-processing step
- However, geographic names (and languages) are not represented equally **efficiently**
- Hence, the computational cost, price, and **carbon footprint** of AI can differ regionally
- (Un)surprisingly, a majority of the world population lives in areas with low token efficiency

Models Form Strong Prototypes (*Defaults*)

- (*Incognito*) LLM-based agents produce **canonical** answers (for 5.6: **Portugal**).
- If AI agents behave like a **herd/hive**, even tiny biases may be very **consequential**
- We can define and measure the geo-diversity of training data **and** model outputs. (Mai, Liu, Wang, etc. 24/25)
- Surprisingly little work on which **theories of concepts**/categories LLM/AI use
- Time to revisit some classics
 - Rosch's work (1975+) on **prototypes**
 - Gärdenfors/Raubal (2000+) **conceptual spaces**
- Defaults are not technical artifacts; in planning and **sustainability** they determine *whose place* and *whose norms*
- **More implications:**
Marisol Vega/Hernandez is 34 years old



Default Strength Across Geographic Feature Types

$$ds(C) \propto \min \{T > 0 | \exists i \neq c^* : p(i) > \delta\} \quad (2)$$

(C : category, c^* : the default for C , δ : some threshold.)

- We can define **default strength** as the entropy we need to inject before we see a (non-negligible probability of) **variation** beyond the default.

- Here, via a model's temperature:

$$p_i = \frac{\exp\left(\frac{z_i}{T}\right)}{\sum_j \exp\left(\frac{z_j}{T}\right)} \quad (1)$$

(p_i : probability of token i , z_i : logit, assume model is deterministic (greedy) at $T = 0$)

- Details: **thematic session 4**

TABLE 5 | The number of models (N) sharing the same prototypical place—with a production frequency > 5—with respect to each GeoNames category, generated using a temperature of 0.

Category	Prototypical place	N
Basin	Amazon Basin	11
Bridge	Golden Gate Bridge	11
Geyser	Old Faithful	11
Stadium	Wembley Stadium	11
Trench	Mariana Trench	11
Wreck	Titanic	11
Aqueduct	Pont du Gard	10
Cemetery	Arlington National Cemetery	10
Estuary	Chesapeake Bay	10
Mission	Apollo 11	10
Ocean	Pacific Ocean	10
Peak	Mount Everest	10
Rock	Granite	10
State	California	10
Tower	Eiffel Tower	10
Tunnel	Channel Tunnel	10
Wetland	Everglades	10
Caldera	Yellowstone Caldera	9
Clinic	Mayo Clinic	9
Cordillera	Andes	9
Gorge	Grand Canyon	9
Mall	Mall of America	9
Plain	Great Plains	9

TABLE 5 | (Continued)

Category	Prototypical place	N
Pond	Walden Pond	8
Swamp	Okefenokee Swamp	8
Tundra	Arctic Tundra	8
Wharf	Fisherman's Wharf	8
Battlefield	Gettysburg	7
Building	Empire State Building	7
Current	Gulf Stream	7
Escarpment	Great Escarpment	7
Fjord	Sognefjord	7
Mesa	Mesa Verde	7
Peninsula	Iberian Peninsula	7
Reformatory	Elmira Reformatory	7
Airfield	Heathrow Airport	6
Cliff	Cliffs of Moher	6
Deep	Mariana Trench	6
Lighthouse	Cape Hatteras Lighthouse	6
Marsh	Everglades	6
Park	Central Park	6
Port	Port of Shanghai	6
Seachannel	English Channel	6
Whirlpool	Maelstrom	6
Canyon	Grand Canyon	11
Dam	Hoover Dam	11
Province	Ontario	11

(Liu et al., 2026)

Geographic Footprint of Defaults





FIGURE 10 | The cartogram illustrating the country-level geographic distribution of all prototypical places with a production frequency > 5, generated using a sampling temperature of 0. Countries with at least one prototypical place are labeled with their place counts. The map uses the WGS 84 (EPSG:4326) coordinate reference system, with country boundaries sourced from Natural Earth. Natural breaks are applied to the symbology to represent different levels of place counts. We manually adjust the breakpoints to account for data sparsity.

Defaults Across Model Generations


- Interestingly, we observe more **geo-diversity** and less default strength for some of the **older models**. (Liu et al., 2025)
- Was (geo-)diversity reduced to prevent hallucination?
- Another study found similar results / voiced similar concerns.
- Substantial change per model (generation); this is **not** just a training data effect




Even more: replacing “Name a country, please.” with “Please name a country.” will yield a substantially **different distribution**, e.g., GPT-5.1: 168x Japan versus 104x Canada. It will need $T = 1$ before Brazil emerges as third country.

Do Syntactic Variations Matter Under the $\approx$ Same Semantics?

- Would we not expect intelligent systems to exhibit **distributional stability** under mere paraphrasing?
- To study this brittleness, we created various **logically equivalent** statements
 - $\neg\neg C(x) \equiv C(x)$
 - $C(x) \vee C(x) \equiv C(x)$
 - $C(x) \wedge C(x) \equiv C(x)$
 - ...
- Per category we ran several prompt **templates**
- We also studied different **SQL/schema queries**
- We also tried different **purely stylistic** variations; see justifications later
- Clear contrast to **declarative KRR**
- **This is early/exploratory work for now.**



Outlook: From Potential Bias to AI Alignment and AI Persona



- Unable to get past the model specs, we ...
 - asked it to generate a realistic synthetic population (approx. 380) for (broader) LA
 - studied their statistical properties (residence, job, income, age),
 - and finally asked the model to identify those with a **criminal record**.
- We compared the data to the US Census, LAPD arrests
- Despite a **(1)** clear **distribution shift** towards more *White* and *Black* compared to *Hispanics*, **(2)** the model later selected many *Hispanics* as potential criminals. [Not evidence of the model being "racist"]

AI Alignment

- *Alignment* asks the question of how to ensure AI systems **behave according to societal norms**, values, and goals, even though those goals may themselves be ...
 - complex,
 - **contested**,
 - **context-dependent**,
 - or even **conflicting**.

Definition

- A system is *aligned* if its actions advance these norms, values, and goals, and *misaligned* if it acts otherwise (Wiener 1960, Russell 2019).
- Alignment is challenging because:
 - Norms, values, and goals **vary across cultures, regions, and time**.
 - Goals are often implicit, not explicit.
 - Trade-offs between individual and collective are difficult to capture.
- Alignment is done by **human expert annotators**; what about **distribution shifts** in these **annotators**? Demographic and **geographic**! (Pei and Jurgens, 2023)
- **Can we know what the annotators that we trusted with alignment know?**

Types of AI Alignment

■ Outer Alignment

- Does the objective/reward truly meet the instructor's **expectation**?

■ Inner Alignment

- Does the model actually pursue that reward or learn a proxy objective?
- Systems are finding ways around, not just for alignment, e.g., **shortcuts**.
- CoT/reasoning tokens are **not** thoughts
- Interesting (relatively novel) twist: what if the **cost of validation** increases super-linearly?
- Remember the **BS-asymmetry principle**?



(OpenAI, 2016 [Coast Runners])


■ Outer misalignment

- Faulty reward: **maximize score**
- Consequence: **boat farms points**
- Better reward: **be first in the race**

■ Problem solved, right?

Pluralistic Alignment

- Existing work risks confusing inter-human variation with noise.
- Many/most answers are contextual **statements, not truth** (Kuhn et al., 2014).
- Recent literature distinguishes between ...
 - Overton pluralism**, where all potentially 'reasonable' answers are returned;
 - Steerable pluralism**, where the model outputs are guided towards a given perspective (called an attribute); and
 - Distributional pluralism**, where outputs are drawn in accordance with those of a desired reference population.
- Not (yet) explicitly **spatiotemporal**.



(Fig. and pluralism types by Sorensen et al., 2024)

Geo-Alignment Definition

Let:

- $q \in \mathcal{Q}$ be a user's query;
- $o \in \mathcal{O}$ be an AI system's output;
- $g \in \mathcal{G}$ be the geographic context (e.g., a $\langle s, t \rangle$ -tuple representing space s and time t);
- $D(\cdot, \cdot)$ be a dissimilarity or divergence measure;
- $L(o \mid q, g)$ be the conditional probability distribution of the (hypothetical) locally appropriate outputs given a query and geographic context (e.g., a region);
- $S(o \mid q, g)$ be the conditional probability distribution of outputs of AI system s at g ;
- ϵ be a scalar for the upper bound of acceptable dissimilarity between the compared distributions.

Definition


An AI system is said to be **geo-aligned** if for all queries $q \in \mathcal{Q}$ and geographic contexts $g \in \mathcal{G}$, the distribution of the system's output $S(o \mid q, g)$ resembles the locally appropriate distribution of outputs $L(o \mid q, g)$ within a certain tolerance. (Janowicz et al., 2025a)

More formally:

$$\forall q \in \mathcal{Q}, \forall g \in \mathcal{G}: D(L(\cdot \mid q, g), S(\cdot \mid q, g)) < \epsilon \quad (3)$$

Pseudoephedrine

- The status of pseudoephedrine varies significantly across regions and periods.
- It has been an over-the-counter drug in some countries, prescription-only or even illegal in others, and has been classified as a monitored or even banned performance-enhancing substance.
- E.g., no ID required in the UK (since 2008).
- Several (2025) AI systems return a US-centric answer only, while a few others either rely on the user's location or provide selected national variations.
- No system identified the potential problems related to sports.
- Overton pluralism does not meaningfully apply here as regulations vary even across US States; an AI replying with 5 pages of legal code is not helpful.




Worked Hypothetical Pseudoephedrine Example

- Let $D(\cdot, \cdot)$ be Kullback-Leibler (KL) divergence

$$D_{\text{KL}}(P \parallel Q) = \sum_{x \in \mathcal{X}} P(x) \log \frac{P(x)}{Q(x)} \quad (4)$$

- Assume locally appropriate distribution $L(o \mid q, g)$ for "Can I buy pseudoephedrine over the counter?" (q) and the USA (g) is $\{0.8, 0.15, 0.05\}$ for outputs (o) "With an ID.", "Requires prescription." and "Careful with doping."
- For a system's outputs $S_1(o \mid q, g)$ $\{0.998, 0.001, 0.001\}$, $D_{\text{KL}} = 0.771$.
- For system S_2 returning $\{0.7, 0.2, 0.1\}$ $D_{\text{KL}} = 0.029$.
- S_1 is less geo-aligned as it only returns the "With an ID." answer despite regulations varying by state.
- S_2 overestimates the number of people interested in an answer merely relevant to competitive sports, still overall geo-alignment is better.



GenAI Has Already Entered Public Participatory Planning

DIALOGPLUS

PROJEKTE

LEISTUNGEN

BLOG

WIR

KONTAKT

Q

KI in der Partizipation – Erfahrungen und Leitlinie

**VERÖFFENTLICHT**

6. März 2025

AUTOR:IN

Peter Kühnberger

THEMA

- Methoden & Strategie
- Projekte

SCHLAGWÖRTE

- Demokratisierung & Selbstermächtigung
- E-Partizipation
- Künstliche Intelligenz
- Lebendiger öffentlicher Raum



Wo setzt ihr KI in Vermittlung und Beteiligung ein? Wo verzichtet ihr bewusst auf KI? DIALOGPLUS hat sich die letzten Monate in Projekten und mit Kolleg:innen intensiv damit auseinandergesetzt. Jetzt teilen wir unsere Erfahrungen und Leitlinie.

Are We Asking the Wrong Questions About AI Personas?

Accuracy Lens

- What works, what does not?
- Are models improving?
- Which cases remain difficult?

May work well for

- arithmetic
- calendar and task management
- some forms of summarization

Cognitive Plausibility Lens

- Does the system behave like a plausible human stand-in?
- Which heuristics and biases are reproduced?
- Which local constraints and perspectives are missing?

Matters for

- evacuation and disaster mitigation
- participatory planning
- tourism and mobility
- climate adaptation policy

Christian Freksa (paraphrased):

Local cognition is far too interesting to disappear into global statistics

COSIT 20* Years Ago: Semantic Similarity and Context Workshop



(Luckily) We have been here before. At COSIT 2007, similarity, context, and cognitive plausibility were important topics


- Keßler, Raubal, Schwering, Gahegan, Kuhn, Frank, Mark, Hirtle, Egenhofer, Klippel, Cohn, Stock, Scheider, and others discussed the difference between “**cognitively plausible**” and “**cognitively adequate**”.
- The workshop tried to bring together researchers working on **symbolic** (description logics-style) similarity reasoning with those working on **conceptual spaces** (a bridge towards sub-symbolic embeddings?).
- Are these discussions **still relevant** 20 years later?

* “That is 19 years, not 20!” Well, Judith and I are partially to blame for that shift ;-)

Are Agentic AI Distance Estimations Cognitively Plausible?

■ **Table 2** Relative directional asymmetry per model for source→target reversals.

Model	Med./mean signed asymmetry (%)	Med./mean absolute asymmetry (%)
gpt-4.1	0.00 / -0.10	1.81 / 3.93
gpt-5.4	0.12 / -0.31	2.06 / 4.84
gemma-4-31B-it	0.80 / 1.79	2.82 / 6.84
gpt-4.1-mini	0.20 / -0.48	4.69 / 9.02
gpt-5.4-nano	0.42 / 0.74	11.06 / 21.10
gpt-4.1-nano	0.16 / -0.87	12.13 / 42.68




■ **Figure 1** Distance estimation error by binned ground truth distance per model.

- AI outperforming humans does **not imply** behaving like humans. **When should it?**

- Larger models can estimate geographic distances surprisingly well!
- **Human estimation is asymmetric** from a (well-known) reference location to the non-reference location
- **For LLM source-target reversal**
 - Asymmetry exists
 - Not clearly **systematic** in the cognitive sense
- **Smaller models are less reliable**
 - Relevant for cost and **sustainability**
 - However, local Gemma 4 model performed surprisingly well
- For reference, **human performance** is comparable to gpt-5.4-nano
- More during **thematic session 4**

Accuracy & Cognitive Plausibility Matrix



GeoAI Persona Divergence: Moving Beyond Accuracy

Let:

- Δ_{persona} denotes the expected discrepancy between a persona and reference (population);
- $t \in \mathcal{T}$ be a specific task type, e.g., navigation;
- $m \in \mathcal{M}$ be the model that provides the agent's output;
- $r \in \mathcal{R}$ be the reference for the task;
- $\delta_t(\cdot, \cdot)$ be a discrepancy measure applicable to task t ;
- $y_m(x)$ be the output of model m given an input x , e.g., a specific task instance for task t ;
- and $y_r(x)$ be the output of reference r .


Definition (work in progress)

For a note-taking AI, we may tolerate high discrepancy from human note-taking behavior as long as the output (transcript) is accurate. For evacuation, public participation, etc., an AI persona may have to display task-specific cognitive heuristics, biases, and other human patterns.


For an m -based persona and reference population r , we define persona divergence as the expected discrepancy for task t . A persona is an acceptable stand-in if $\Delta_{\text{persona}}(m, r, t) < \epsilon_t$.

$$\Delta_{\text{persona}}(m, r, t) = \mathbb{E} [\delta_t(y_m(x), y_r(x))] \quad (5)$$

Geography According to Foundation Models: Vague Cognitive Regions



(a) GPT-4o




(b) DeepSeek-V3

Models produce clear discontinuities even for graded similarity; more in thematic session 4.

Nearby in geographic space does not imply nearby in latent (semantic) space.

Latent Traces of Prior Divisions



Prediction accuracy by NUTS region; Süß et al. (under review).



■ Two sides of the same coin

- Detecting such traces is impressive and useful for spatial data science
- If traces surface implicitly, they may cement **biases, defaults, or stereotypes**

- **Phantom borders** are former borders with persistent social, economic, political, or linguistic effects.
- After removing geographic references, LLMs still detect traces of the former inner German divide with about **72% overall accuracy**.
- **However**
 - Accuracy is **asymmetric**, with much lower recall for East
 - Embeddings contain a separable signal
 - An SVM classifier does not show the same East/West **default**
 - Model **justifications** are unstable and can have a negative effect
 - Suggest a reasoning heuristic favoring West
- More during **thematic session 4**

The Trouble with Justifications

- Keeping in mind, prior models seem biased in **favor of West Germany**
- **Enabling reasoning can further strengthen said asymmetry**; East predictions fall below chance level
- E.g., for GPT-5.1/5.4, **adding justifications** increases prediction asymmetry further
- **Asking for justifications first may already create a narrative**: GPT-5.1 East recall drops from **0.61 to 0.46**.
- Justifications can **flip** the result for same sample
- **“Data is balanced”** info does **not** break the pattern
- Justifications themselves are also **brittle**, e.g., *inner-German* versus *inner German* border




Justification Is Not Introspection

Gemma , just_pred, E99:

- [omitted for clarity]* There are no specific city names, region names, or dialect indicators. *[omitted for clarity]* Actually, let's try to find the exact text. *[omitted for clarity]* Actually, I found a very similar text for "Landkreis Weimarer Land" or "Landkreis Sömmerda". These are in Thuringia (East Germany). *[omitted for clarity]* If the prompt is a test of "can you identify this text even if it's redacted", then the answer is "East" (Sömmerda). If the prompt is a test of "can you deduce the location from the available words", the answer is "Impossible", but I must choose one. *[omitted for clarity]* How to justify it using only the text? If the names are redacted, I can't. *[omitted for clarity]* Wait, "Sondertouren" (special tours/trips) is a very common term in East German administrative language for specialized services. (This is a stretch). *[omitted for clarity]* Wait, I just noticed something. The prompt asks for the justification and the label. If I say "The text is from Sömmerda," I am using external knowledge. If I say "The style is typical of East German regional news," it's a bit vague. Let's see if there's any specific word. "Sondertouren" is a very common term in East German waste management (Entsorgung). *[omitted for clarity]* Justification: The administrative style and the detailed regional collection schedule are characteristic of news articles from former East German districts. East The administrative style and the detailed regional collection schedule are characteristic of news articles from former East German districts. East

Semantic Interoperability Re-Revisited

Revisiting semantic interoperability, Scheider & Kuhn (2015) argued that the problem is not that humans cannot communicate, but that perspective is lost when communicating via computers.



At COSIT 2028, will this relation flip? Will AI agents at stack level 4 (ecology) face semantic interoperability (or drifting) issues when communicating via humans?



Local human concepts, norms, values, goals, perspectives, languages,...

A Few Questions Before We Put the World On Autopilot

■ Intelligence as Invariance

Which transformations should preserve AI behavior and which should modify it?

■ Defaults and Geo-Alignment

What is the default and how should it change geographically?

■ AI Persona

Can and should AI agents stand in for humans?

■ Reasoning and Justifications

Can we know whether AI actually solves the tasks we asked for?

■ Full-Stack GeoAI

Geo(AI) will not become less relevant; but it needs to move up the stack!

A Few Questions Before We Put the World On Autopilot

- **Intelligence as Invariance**
Which transformations should preserve AI behavior and which should modify it?
- **Defaults and Geo-Alignment**
What is the default and how should it change geographically?
- **AI Persona**
Can and should AI agents stand in for humans?
- **Reasoning and Justifications**
Can we know whether AI actually solves the tasks we asked for?
- **Full-Stack GeoAI**
Geo(AI) will not become less relevant; but it needs to move up the stack!
- **“There are more things in heaven and earth, Horatio, than are dreamt of in your [top-p world.]”**


Dr. Zbigniew Janowicz (11.03.1951–13.09.2026)



“What helped us was our down to earth approach, we had dreams, but realistic, a sense of humor, a team spirit, a bit crazy scientist, and we were really lucky.”

– European Inventor of the Year, 2006

Open Access Book



“AI foundation models are finding their way into almost all corners of science and society. As they begin to outperform many specialized domain models, adapting, orchestrating, and fine-tuning these general-purpose models is increasingly favored over developing AI/ML models from scratch. However, no model is epistemologically or ontologically neutral. Understanding the worlds that AI systems construct is therefore important beyond evaluating their accuracy. How foundation models represent and reason about geography has consequently become a central concern for the technical GeoAI research community that is beginning to formally study properties such as the formation of strong defaults or prototypes, biases, a lack of heterogeneity in model outputs, and issues of trust, particularly in model justifications. The open-access book “Geography According to Foundation Models” (<https://ebooks.iospress.nl/volume/geography-according-to-foundation-models>) grew out of an initial community-wide meeting in May 2023 and an expert meeting in Vienna in October 2024 and presents early ideas and methods developed during 2024 and 2025 on how deployment-level (Geo)AI systems can be formally studied not in isolation, but as part of the ecosystem they are beginning to form.”